

Lecture 4: persistence with mongodb

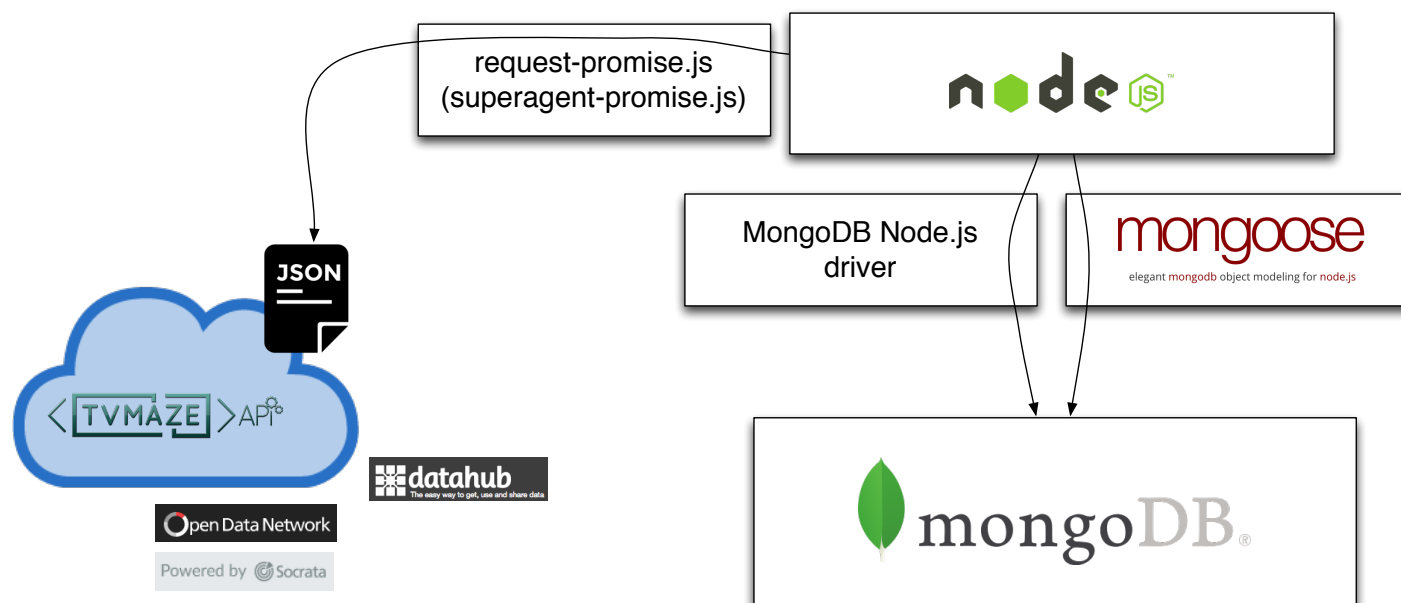
Olivier Liechti
TWEB

heig-vd

Haute Ecole d'Ingénierie et de Gestion
du Canton de Vaud

Webcasts

- 1 use one of the "promisified" npm modules to fetch JSON data from the Web



- 2 store the JSON documents in MongoDB collections and run queries

Tasks

1. Prepare the environment

- 1.1. Start a MongoDB server in a Docker container
- 1.2. Start a MongoDB shell in another Docker container
- 1.3. Run basic commands (create a document, find all documents)

2. Get data from the Web

- 2.1. Find JSON data source
- 2.2. Pick a npm module to submit HTTP requests
- 2.3. Fetch data

3. Access MongoDB with the mongo driver

- 3.1. Handle the connection with the MongoDB server
- 3.2. Explore the driver API
- 3.3. Interact with collections and documents

4. Integrate all the pieces with Promises

- 3.1. Callbacks vs promises in the APIs
- 3.2. Promise chains
- 3.3. Passing data and returning results

5. Access MongoDB with Mongoose

- 4.1. Select and study a template
- 4.2. Discover jquery datatables
- 4.3. Integrate the template in the project

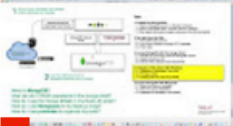
What is MongoDB?

How can do I CRUD operations in the mongo shell?

How do I use the mongo **driver** in my Node.JS script?

How do I use **Mongoose** in my Node.js script?

How do I use **promises** to organize my code?

11		Bootcamp 4.1: Intro aux webcasts "MongoDB" by oliechti	2:48
12		Bootcamp 4.2: prise en main de MongoDB by oliechti	20:30
13		Bootcamp 4.3 (a): identification de la source de données JSON by oliechti	5:40
14		Bootcamp 4.3 (b): utilisation de request-promise pour interroger l'API REST by oliechti	15:12
15		Bootcamp 4.4: utilisation du driver node.js MongoDB by oliechti	21:01
16		Bootcamp 4.4: implémentation de la chaîne de promesses by oliechti	38:46

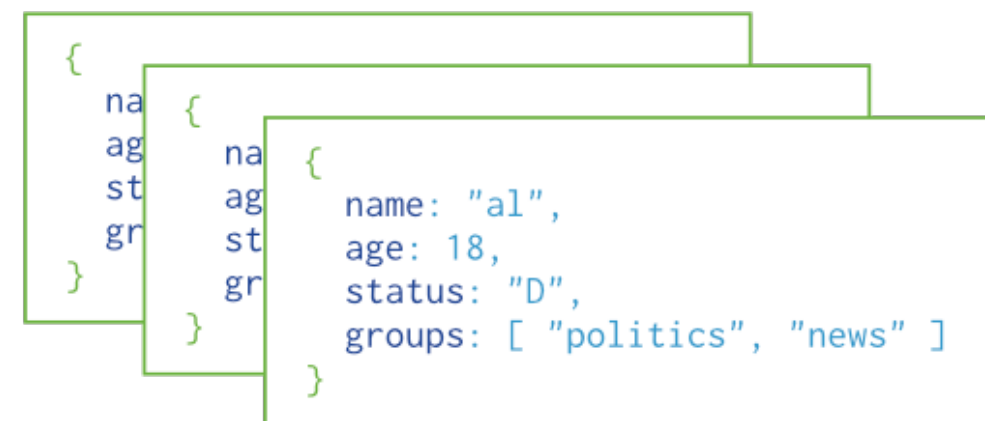


Document-oriented NoSQL Database

- MongoDB is one of the most popular **NoSQL** databases (and one of the first to have been categorized as such).
- It is a **schema-less document-oriented** database:
 - The data store is made of several **collections**.
 - Every collection contains a set of **documents**, which you can think of as JSON objects.
 - The **structure of documents is not defined *a priori*** and is not enforced. This means that a collection can contain documents that have different fields.

```
{  
  name: "sue",  
  age: 26,  
  status: "A",  
  groups: [ "news", "sports" ]  
}
```

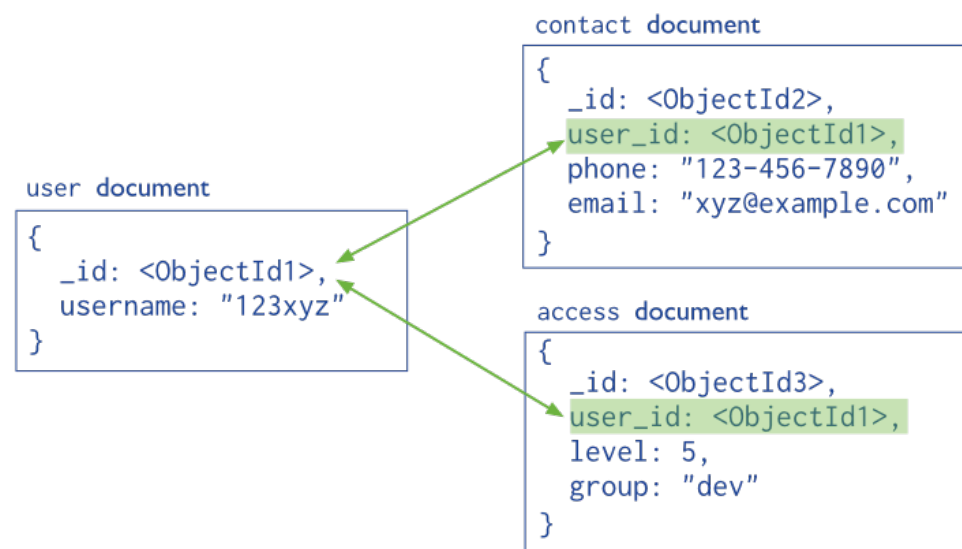
← field: value
← field: value
← field: value
← field: value



Collection

Data modeling

- Creating a data model with MongoDB does not have to follow the rules that apply for relational databases. Often, they should not.
- Consider these questions: is this a **composition** relationship (**containment**)? Is this "**aggregate**" of documents often used at the same time (i.e. can we reduce chattiness)? Would embedding lead to "a lot" of data duplication?



Normalized data model
(references)



Embedded data model
(sub-documents)

One-to-one relationships

2 documents (requires 2 queries
to get all of the person data)

```
{
  _id: "joe",
  name: "Joe Bookreader"
}

{
  patron_id: "joe",
  street: "123 Fake Street",
  city: "Faketon",
  state: "MA",
  zip: "12345"
}
```

Normalized data model
(references)

```
{
  _id: "joe",
  name: "Joe Bookreader",
  address: {
    street: "123 Fake Street",
    city: "Faketon",
    state: "MA",
    zip: "12345"
  }
}
```

1 single, aggregate document
(in this case, it is a better choice)

Embedded data model
(sub-documents)

<https://docs.mongodb.org/manual/tutorial/model-embedded-one-to-one-relationships-between-documents/>

One-to-many relationships

```
{
  _id: "joe",
  name: "Joe Bookreader"
}

{
  patron_id: "joe",
  street: "123 Fake Street",
  city: "Faketon",
  state: "MA",
  zip: "12345"
}

{
  patron_id: "joe",
  street: "1 Some Other Street",
  city: "Boston",
  state: "MA",
  zip: "12345"
}
```

```
{
  _id: "joe",
  name: "Joe Bookreader",
  addresses: [
    {
      street: "123 Fake Street",
      city: "Faketon",
      state: "MA",
      zip: "12345"
    },
    {
      street: "1 Some Other Street",
      city: "Boston",
      state: "MA",
      zip: "12345"
    }
  ]
}
```



MongoDB document can have
an arbitrary structure, including
arrays

One-to-many relationships

heig-vd

Haute Ecole d'Ingénierie et de Gestion
du Canton de Vaud

ok if if have few books per publisher

```
{
  title: "MongoDB: The Definitive Guide",
  author: [ "Kristina Chodorow", "Mike Dirolf" ],
  published_date: ISODate("2010-09-24"),
  pages: 216,
  language: "English",
  publisher: {
    name: "O'Reilly Media",
    founded: 1980,
    location: "CA"
  }
}

{
  title: "50 Tips and Tricks for MongoDB Developer",
  author: "Kristina Chodorow",
  published_date: ISODate("2011-05-06"),
  pages: 68,
  language: "English",
  publisher: {
    name: "O'Reilly Media",
    founded: 1980,
    location: "CA"
  }
}
```

A red arrow points from the word "duplication" to the publisher object in the first document. Another red arrow points from the word "duplication" to the publisher object in the second document.

duplication

```
{
  name: "O'Reilly Media",
  founded: 1980,
  location: "CA",
  books: [123456789, 234567890, ...]
}

{
  _id: 123456789,
  title: "MongoDB: The Definitive Guide",
  author: [ "Kristina Chodorow", "Mike Dirolf" ],
  published_date: ISODate("2010-09-24"),
  pages: 216,
  language: "English"
}

{
  _id: 234567890,
  title: "50 Tips and Tricks for MongoDB Developer",
  author: "Kristina Chodorow",
  published_date: ISODate("2011-05-06"),
  pages: 68,
  language: "English"
}
```

A red arrow points from the text "ok if if have few books per publisher" to the 'books' array in the publisher document.

```
{
  _id: "oreilly",
  name: "O'Reilly Media",
  founded: 1980,
  location: "CA"
}

{
  _id: 123456789,
  title: "MongoDB: The Definitive Guide",
  author: [ "Kristina Chodorow", "Mike Dirolf" ],
  published_date: ISODate("2010-09-24"),
  pages: 216,
  language: "English",
  publisher_id: "oreilly"
}

{
  _id: 234567890,
  title: "50 Tips and Tricks for MongoDB Developer",
  author: "Kristina Chodorow",
  published_date: ISODate("2011-05-06"),
  pages: 68,
  language: "English",
  publisher_id: "oreilly"
}
```

A red arrow points from the text "better if you have many books per publisher" to the 'publisher_id' field in the second document.

better if you have many books per publisher

<https://docs.mongodb.org/manual/tutorial/model-embedded-one-to-many-relationships-between-documents/>

Insert data in MongoDB

- To insert data in MongoDB, you simply have to provide a **JSON document** (with an **arbitrary structure**).
- The documents in the collection do not have to all have the same structure (this is why we talk about a **schemaless** database).

A document is always inserted
in a specific collection.

MongoDB will assign a unique
ID in the **_id** field for the new
document.

Collection
↓
`db.users.insert(`

Document
↓
{
 name: "sue",
 age: 26,
 status: "A",
 groups: ["news", "sports"]
}

)

Document

```
{  
  name: "sue",  
  age: 26,  
  status: "A",  
  groups: [ "news", "sports" ]  
}
```

insert

Collection

{ name: "al", age: 18, ... }
{ name: "lee", age: 28, ... }
{ name: "jan", age: 21, ... }
{ name: "kai", age: 38, ... }
{ name: "sam", age: 18, ... }
{ name: "mel", age: 38, ... }
{ name: "ryan", age: 31, ... }
{ name: "sue", age: 26, ... }

users

Insert data in MongoDB

- This is one example. You can also insert multiple documents at the same time, either by passing an array of documents or by performing a bulk operation. If you are dealing with many documents, this is important for performance reasons.

```
db.inventory.insert(  
  {  
    item: "ABC1",  
    details: {  
      model: "14Q3",  
      manufacturer: "XYZ Company"  
    },  
    stock: [ { size: "S", qty: 25 }, { size: "M", qty: 50 } ],  
    category: "clothing"  
  }  
)
```

```
db.inventory.find()
```

```
{ "_id" : ObjectId("53d98f133bb604791249ca99"), "item" : "AB
```

Update and delete data in MongoDB

- When you update or delete documents, you specify which documents are concerned by the operation.
- You do that by specifying update or remove criteria. Specifying "{}" means that you want to apply the operation on all documents of the collection.

```
db.users.update(  
  { age: { $gt: 18 } },  
  { $set: { status: "A" } },  
  { multi: true }  
)
```

← collection
← update criteria
← update action
← update option

```
db.users.remove(  
  { status: "D" }  
)
```

← collection
← remove criteria

Update data in MongoDB

- By default, update will modify only the first document that matches the selection criteria. You can specify the "multi" options if you want to update all documents that match the criteria.

```
db.inventory.update(  
  { item: "MNO2" },  
  {  
    $set: {  
      category: "apparel",  
      details: { model: "14Q3", manufacturer: "XYZ Company" }  
    },  
    $currentDate: { lastModified: true }  
  },  
)
```

selection criteria

update operators

lastModified will be set to the
current date

```
db.inventory.update(  
  { category: "clothing" },  
  {  
    $set: { category: "apparel" },  
    $currentDate: { lastModified: true }  
  },  
  { multi: true }  
)
```

Update data in MongoDB

heig-vd

Haute Ecole d'Ingénierie et de Gestion
du Canton de Vaud

Update Operators

Fields

Name	Description
<code>\$inc</code>	Increments the value of the field by the specified amount.
<code>\$mul</code>	Multiplies the value of the field by the specified amount.
<code>\$rename</code>	Renames a field.
<code>\$setOnInsert</code>	Sets the value of a field if an update results in an insert of a document. Has no effect on update operations that modify existing documents.
<code>\$set</code>	Sets the value of a field in a document.
<code>\$unset</code>	Removes the specified field from a document.
<code>\$min</code>	Only updates the field if the specified value is less than the existing field value.
<code>\$max</code>	Only updates the field if the specified value is greater than the existing field value.
<code>\$currentDate</code>	Sets the value of a field to current date, either as a Date or a Timestamp.

Array

Operators

Name	Description
<code>\$</code>	Acts as a placeholder to update the first element that matches the query condition in an update.
<code>\$addToSet</code>	Adds elements to an array only if they do not already exist in the set.
<code>\$pop</code>	Removes the first or last item of an array.
<code>\$pullAll</code>	Removes all matching values from an array.
<code>\$pull</code>	Removes all array elements that match a specified query.
<code>\$pushAll</code>	<i>Deprecated.</i> Adds several items to an array.
<code>\$push</code>	Adds an item to an array.

<https://docs.mongodb.org/manual/reference/operator/update/>

Query MongoDB

- MongoDB is one of the most popular **NoSQL** databases (and one of the first to have been categorized as such).

Collection
`db.users.find(` Query Criteria `{ age: { $gt: 18 } }).sort(` Modifier `{age: 1 })`

{ age: 18, ... }
{ age: 28, ... }
{ age: 21, ... }
{ age: 38, ... }
{ age: 18, ... }
{ age: 38, ... }
{ age: 31, ... }

users

Query Criteria

{ age: 28, ... }
{ age: 21, ... }
{ age: 38, ... }
{ age: 38, ... }
{ age: 31, ... }

Modifier

{ age: 21, ... }
{ age: 28, ... }
{ age: 31, ... }
{ age: 38, ... }
{ age: 38, ... }

Results

We optionally sort the documents, limit the number of documents returned, etc.

We define the criteria for which documents should be considered, and which of their fields should be considered (projection)

<https://docs.mongodb.org/manual/core/crud-introduction/>

We look for documents within one collection, within one database

Query MongoDB

```
db.inventory.find()
```

all documents

```
db.inventory.find( { type: "snacks" } )
```

documents, which have a field "type"
with a value of "snacks"

```
db.inventory.find( { type: { $in: [ 'food', 'snacks' ] } } )
```

documents, which have a field "type" with have a value
of "food" or "snacks"

```
db.inventory.find( { type: 'food', price: { $lt: 9.95 } } )
```

documents with a type field equal to "food" AND a price
field with a value less than 9.95

```
db.inventory.find(
  {
    $or: [ { qty: { $gt: 100 } }, { price: { $lt: 9.95 } } ]
  }
)
```

documents where the quantity is
more than 100 OR the price is
less than 9.95

Query MongoDB: arrays

```
{ _id: 5, type: "food", item: "aaa", ratings: [ 5, 8, 9 ] }  
{ _id: 6, type: "food", item: "bbb", ratings: [ 5, 9 ] }  
{ _id: 7, type: "food", item: "ccc", ratings: [ 9, 5, 8 ] }
```

```
db.inventory.find( { ratings: [ 5, 8, 9 ] } )
```

```
{ "_id" : 5, "type" : "food", "item" : "aaa", "ratings" : [ 5, 8,
```

Exact match on the entire array

```
db.inventory.find( { ratings: 5 } )
```

Return documents if the array
contains a specific value

```
{ "_id" : 5, "type" : "food", "item" : "aaa", "ratings" : [ 5, 8, 9 ] }  
{ "_id" : 6, "type" : "food", "item" : "bbb", "ratings" : [ 5, 9 ] }  
{ "_id" : 7, "type" : "food", "item" : "ccc", "ratings" : [ 9, 5, 8 ] }
```

Query MongoDB: arrays

```
{ _id: 5, type: "food", item: "aaa", ratings: [ 5, 8, 9 ] }  
{ _id: 6, type: "food", item: "bbb", ratings: [ 5, 9 ] }  
{ _id: 7, type: "food", item: "ccc", ratings: [ 9, 5, 8 ] }
```

```
db.inventory.find( { ratings: { $elemMatch: { $gt: 5, $lt: 9 } } } )
```

Documents where one element of ratings is at the same time > 5 AND < 9

```
{ "_id" : 5, "type" : "food", "item" : "aaa", "ratings" : [ 5, 8, 9 ] }  
{ "_id" : 7, "type" : "food", "item" : "ccc", "ratings" : [ 9, 5, 8 ] }
```

```
db.inventory.find( { ratings: { $gt: 5, $lt: 9 } } )
```

Documents where there is one element of ratings > 5 and one element < 9

```
{ "_id" : 5, "type" : "food", "item" : "aaa", "ratings" : [ 5, 8, 9 ] }  
{ "_id" : 6, "type" : "food", "item" : "bbb", "ratings" : [ 5, 9 ] }  
{ "_id" : 7, "type" : "food", "item" : "ccc", "ratings" : [ 9, 5, 8 ] }
```

SQL vs MongoDB

heig-vd

Haute Ecole d'Ingénierie et de Gestion
du Canton de Vaud

```
SELECT *  
FROM users  
WHERE status != "A"
```

```
db.users.find(  
  { status: { $ne: "A" } }  
)
```

```
SELECT *  
FROM users  
WHERE status = "A"  
AND age = 50
```

```
db.users.find(  
  { status: "A",  
    age: 50 }  
)
```

```
SELECT *  
FROM users  
WHERE status = "A"  
OR age = 50
```

```
db.users.find(  
  { $or: [ { status: "A" } ,  
           { age: 50 } ] }  
)
```

```
SELECT *  
FROM users  
WHERE age > 25
```

```
db.users.find(  
  { age: { $gt: 25 } }  
)
```

```
SELECT *  
FROM users  
WHERE age < 25
```

```
db.users.find(  
  { age: { $lt: 25 } }  
)
```

```
SELECT *  
FROM users  
WHERE age > 25  
AND age <= 50
```

```
db.users.find(  
  { age: { $gt: 25, $lte: 50 } }  
)
```

```
SELECT *  
FROM users  
WHERE user_id like "%bc%"
```

```
db.users.find( { user_id: /bc/ } )
```

```
SELECT *  
FROM users  
WHERE status = "A"  
ORDER BY user_id ASC
```

```
db.users.find( { status: "A" } ).sort( { user_id: 1 } )
```

```
SELECT *  
FROM users  
WHERE status = "A"  
ORDER BY user_id DESC
```

```
db.users.find( { status: "A" } ).sort( { user_id: -1 } )
```

```
SELECT COUNT(*)  
FROM users
```

```
db.users.count()
```

or

```
db.users.find().count()
```

```
SELECT COUNT(user_id)  
FROM users
```

```
db.users.count( { user_id: { $exists: true } } )
```

or

```
db.users.find( { user_id: { $exists: true } } ).count()
```

```
SELECT COUNT(*)  
FROM users  
WHERE age > 30
```

```
db.users.count( { age: { $gt: 30 } } )
```

or

```
db.users.find( { age: { $gt: 30 } } ).count()
```

```
SELECT DISTINCT(status)  
FROM users
```

```
db.users.distinct( "status" )
```

<https://docs.mongodb.org/manual/reference/sql-comparison/>

Query MongoDB: projections

- When you perform a query, you can specify which fields you are interested in (think about performance)

```
db.inventory.find( { type: 'food' } )
```

Return all fields (no second argument)

```
db.inventory.find( { type: 'food' }, { item: 1, qty: 1 } )
```

We are only interested by the item and qty fields (we will also get _id)

```
db.inventory.find( { type: 'food' }, { item: 1, qty: 1, _id: 0 } )
```

We are only interested by the item and qty fields (we really don't want to get _id)

```
db.inventory.find( { type: 'food' }, { type: 0 } )
```

We want all fields except type

<https://docs.mongodb.org/manual/tutorial/project-fields-from-query-results/>

mongojs

 Star 1,099

A [node.js](#) module for mongodb, that emulates [the official mongodb API](#). It wraps [mongodb-native](#) and is available through [npm](#)

```
npm install mongojs
```

build passing



Accessing MongoDB from Node.js

Accessing mongoDB from Node.js

- In the **Java ecosystem**, it is possible to interact with a RDBMS by using a JDBC driver:
 - The program loads the driver.
 - The program establishes a connection with the DB.
 - The program sends SQL queries to read and/or update the DB.
 - The program manipulates tabular result sets returned by the driver.
- With **Node.js and mongoDB**, the process is similar:
 - There is a **Node.js driver for mongoDB** (in fact, there are several).
 - A Node.js module can connect to a mongoDB server and issue queries to **manipulate collection and documents**.

Example 1: connect and insert

```
var MongoClient = require('mongodb').MongoClient;

MongoClient.connect("mongodb://localhost:27017/exampleDb", function(err, db) {
  if(err) { return console.dir(err); }

  var collection = db.collection('test');

  var doc1 = {'hello':'doc1'};
  var doc2 = {'hello':'doc2'};
  var lotsOfDocs = [{'hello':'doc3'}, {'hello':'doc4'}];

  collection.insert(doc1);
  collection.insert(doc2, {w:1}, function(err, result) {});
  collection.insert(lotsOfDocs, {w:1}, function(err, result) {});

});
```

Example 2: query

```
var MongoClient = require('mongodb').MongoClient;

MongoClient.connect("mongodb://localhost:27017/exampleDb", function(err, db) {
  if(err) { return console.dir(err); }

  var collection = db.collection('test');
  var docs = [{mykey:1}, {mykey:2}, {mykey:3}];
  collection.insert(docs, {w:1}, function(err, result) {

    // beware of memory consumption!
    collection.find().toArray(function(err, items) {});

    // better when many documents are returned
    var stream = collection.find({mykey:{$ne:2}}).stream();
    stream.on("data", function(item) {});
    stream.on("end", function() {});

    // special case when only one document is expected
    collection.findOne({mykey:1}, function(err, item) {});

  });
});
```


Accessing mongoDB from Node.js

- **mongojs** is a very useful npm module, which provides an alternative to the official Node.js driver. Its API is very similar to what you type on the mongo

```
// simple usage for a local db
var db = mongojs('mydb', ['mycollection']);

// the db is on a remote server (the port default to mongo)
var db = mongojs('example.com/mydb', ['mycollection']);

// we can also provide some credentials
var db = mongojs('username:password@example.com/mydb', ['mycollection']);

// connect now, and worry about collections later
var db = mongojs('mydb');
var mycollection = db.collection('mycollection');
```

```
// find everything
db.mycollection.find(function(err, docs) {
  // docs is an array of all the documents in mycollection
});

// find everything, but sort by name
db.mycollection.find().sort({name:1}, function(err, docs) {
  // docs is now a sorted array
});
```

mongojs

Star 1,099

A node.js module for mongodb, that emulates the official mongodb API. It wraps [mongodb-native](#) and is available through [npm](#)

npm install mongojs

build passing

mongoose

elegant **mongodb** object modeling for **node.js**



mongoDB

Object Document Mapping with Mongoose

Mongoose: an ORM for MongoDB

- In the **Java EE ecosystem**, we have seen how the **Java Persistence API** (JPA) specifies a standard way to interact with Object-Relational Mapping (ORM) frameworks.
 - The developer **first** creates an **object-oriented domain model**, by creating Entity classes and using various annotations (@Entity, @Id, @OneToMany, @Table, etc.)
 - He **then** uses an **Entity Manager** to **Create, Read, Update and Delete** objects in the DB.
 - The ORM framework takes care of the details: it **generates the schema** and the **SQL queries**.
- In the Javascript ecosystem, we have similar mechanisms. **With the particular yeoman generator that we use for the project:**
 - The authors have decided not use a relational database, but rather the **mongodb** document-oriented database.
 - They have decided to use **one of the data mapping tools** available for mongodb, namely **mongoose**. Since mongodb is a document-oriented database, it is more appropriate to talk about an Object-Document Mapping tool, rather than an ORM.

Mongoose: an ODM for MongoDB

*“Mongoose provides a **straight-forward**, schema-based solution to **modeling** your application data and includes built-in type **casting**, **validation**, **query** building, business logic hooks and more, out of the box.”*

Schema

“Everything in Mongoose starts with a Schema. Each **schema maps to a MongoDB collection** and defines the shape of the documents within that collection.”

Model

“**Models are fancy constructors** compiled from our Schema definitions.”

Document

“Mongoose documents represent a **one-to-one mapping** to documents as stored in MongoDB. Each document is an **instance of its Model**.”

Example

```
var userSchema = new mongoose.Schema({
  name: {
    first: String,
    last: { type: String, trim: true }
  },
  age: { type: Number, min: 0 }
});
```

```
var PUser = mongoose.model('PowerUsers', userSchema);
```

```
var johndoe = new PUser ({
  name: { first: 'John', last: ' Doe ' },
  age: 25
});
```

```
johndoe.save(function (err) {if (err) console.log ('Error on save!')});
```

document

Example: query

we can chain conditions

```
Person
.find({ occupation: /host/ })
.where('name.last').equals('Ghost')
.where('age').gt(17).lt(66)
.where('likes').in(['vaporizing', 'talking'])
.limit(10)
.sort('-occupation')
.select('name occupation')
.exec(callback);
```

we are interested in only some of the fields

we only want to get at most 10 documents

Yo express!

yo angular

```
// Example model

var mongoose = require('mongoose'),
    Schema = mongoose.Schema;

var ArticleSchema = new Schema({
  title: String,
  url: String,
  text: String
});

ArticleSchema.virtual('date')
  .get(function(){
    return this._id.getTimestamp();
  });

mongoose.model('Article', ArticleSchema);
```

app/models/article.js

Yo express!

yo angular

```
var express = require('express'),
    router = express.Router(),
    mongoose = require('mongoose'),
    Article = mongoose.model('Article');

module.exports = function (app) {
  app.use('/', router);
};

router.get('/', function (req, res, next) {
  Article.find(function (err, articles) {
    if (err) return next(err);
    res.render('index', {
      title: 'Generator-Express MVC',
      articles: articles
    });
  });
});
```

app/controllers/home.js